


TL;DR

The “Too Long; Didn’t Read” Privacy Extension

“Nobody reads the fine print. We do it for you.”

 Team Nexus — Devasya Gupta (202302626010025)
— Meet Patel (202302626010053)

GLS University, Ahmedabad • Challenge by Everest Instruments



THE PROBLEM

You “Agree” to things you never read.

WORK DAYS PER YEAR

76

estimated time (Carnegie Mellon’s McDonald & Cranor study) to read every privacy policy the average person agrees to in a year.

This is a trust and transparency crisis at internet scale.

Every install, every signup, every site asks you to accept a Terms of Service or Privacy Policy — dense, thousand-word legal documents almost no one actually reads before clicking “Agree.” Blind acceptance means users unknowingly sign away data rights, consent to AI training on their content, or accept clauses that quietly harm their privacy.



Data sharing & selling permissions

Privacy policies may grant companies permission to share or sell personal data.



AI training on your content

Your words, photos, and files used to train models — often opt-out only.



Indefinite data retention

No clear deadline for when your information is actually deleted.



Zero visibility, zero leverage

Users have no easy way to see what they’ve actually agreed to.

This isn't a guess — it's a documented problem.

Four data points that shaped this project, cited from primary sources — not generated.



9%

of U.S. adults say they always read a privacy policy before agreeing — 56% click “agree” without reading at all.

Pew Research Center, 2023



76 days

of full-time reading a year would be needed to actually read every privacy policy the average person encounters.

McDonald & Cranor, Carnegie Mellon University



16%

of Indian consumers understand the country's own Digital Personal Data Protection Act — despite it now being in force.

DPDP Rules, 2025 (MeitY) — industry compliance research, 2026



Human-only

the closest existing tool (ToS;DR) relies entirely on volunteer review — it cannot scale to the web's millions of policies.

Terms of Service; Didn't Read, est. 2012



TL;DR reads the fine print instantly.

A lightweight browser extension that demystifies any Terms of Service or Privacy Policy the moment you land on it.



01

Detect

Scrapes the on-page legal text the instant a ToS / Privacy Policy page loads.



02

Analyze

Sends the text to an LLM API with a prompt tuned to hunt for risky clauses.



03

Summarize

Returns a clean, risk-ranked 3-bullet summary of what actually matters.



04


Reveal

Displays it instantly in a popup — no tab switching, no legal jargon.

Zero copy-pasting. Zero extra tabs. One click to clarity.



From page load to risk-ranked summary in seconds.



Popup output is color-coded by risk level — so severity is understood at a glance, not buried in a paragraph.

WHAT IT LOOKS LIKE

Not just another "summarize this page" tool.

✓ Risk-first, not generic

Prompted specifically to hunt for data monetization, AI training rights, retention terms, and third-party sharing.

✓ Zero friction


Works passively in the background. One click. No copy-pasting text into another tool.

✓ Fast by design

Built to feel instant — a quick glance, not a heavyweight analysis session.

✓ Works everywhere

Universal across any site with a ToS or Privacy Policy — not tied to a single platform.




Actual popup UI concept — risk-ranked, color-coded, and rendered in under two seconds.

Why TL;DR — and not the alternatives?

Four ways people currently handle this problem, compared side by side.

	Reading it yourself	Copy-paste to ChatGPT	ToS;DR (existing tool)	TL;DR (us)
Speed	✗	⊖	✓	✓
Coverage	✓	✓	⊖	✓
Risk-focused analysis	✗	✗	⊖	✓
Zero effort / friction	✗	✗	⊖	✓
Always up to date	✓	✓	✗	✓

Lightweight, stateless, privacy-respecting by design.





The risk-ranking algorithm, in pseudocode.

A simplified view of the detect → classify → summarize → rank pipeline running behind the popup.

```
risk-engine.pseudo

FUNCTION analyzePage(pageContent):
    IF NOT isLegalPage(pageContent): RETURN null

    legalText = extractText(pageContent)
    IF length(legalText) < MIN_THRESHOLD:
        RETURN null

    // pre-filter with lightweight ML classifier
    riskyChunks = classifyClauses(legalText)

    // build a risk-focused prompt for the LLM
    prompt = buildRiskPrompt(legalText, riskyChunks)
    response = callLLM(prompt)

    summary = parseStructured(response)
    FOR clause IN summary.clauses:
        clause.riskLevel = rankRisk(clause.category)

    RETURN sortByRisk(summary.clauses)[0:3]
```



Assumptions, challenges & open questions.




ASSUMPTIONS

- The extension requests only the browser permissions required to read the current page.
- Target ToS / Privacy pages render as parsable text, not scanned images.
- The chosen LLM API can reliably parse policy language across formats and phrasing.



CHALLENGES

- Heavy JS-rendered sites may delay when legal text actually becomes scrapeable.
- Separating genuinely risky clauses from routine legal boilerplate at scale.
- Keeping the detect → summarize → render loop feeling instant under real latency.



OPEN QUESTIONS

- Which LLM provider best fits the hackathon's API credit / rate-limit constraints?
- Should the MVP prioritize breadth (many sites) or depth (accuracy on top platforms)?
- How much clause detail should the popup show before it feels like "too long" again?

Shifting power back to the user.

In an increasingly consent-fatigued internet, TL;DR gives the individual real leverage back.



Constant exposure

Every internet user hits ToS agreements dozens of times a month — and reads virtually none of them.



A growing concern

Public trust in AI companies training on personal data without clear consent is eroding fast. TL;DR helps users catch it.




Informed, not just compliant

Users move from blindly clicking “Agree” to actually understanding what they’re signing away.



From hackathon MVP to a privacy-native browser habit.



MVP — HACKATHON BUILD

- Core scrape → LLM summarize → popup display
- Working on major platforms



NEXT

- Risk-scoring system with color-coded severity
- Summary history / dashboard
- Support for mobile browsers



FUTURE

- Native alerts before clicking "Agree" on a risky clause
- Aggregated policy-change tracking & notifications

Team Nexus.



Devasya Gupta

202302626010025



Meet Patel

202302626010053



Nobody reads the fine print. We do it for you.

Built by Team Nexus — ready to build the full MVP in the next round.



Team Nexus — Devasya Gupta • Meet Patel

TL;DR • GLS Nexus Hackathon 2026 • Challenge by Everest Instruments